

DATA130008 Introduction to Artificial Intelligence

复旦大学大数据学院
School of Data Science, Fudan University

魏忠钰

Search I

September 14th, 2021


Outline

- Search Problems

Search Problems

- **A search problem consists of:**
 - **A state space**
 - State describes the status of a piece of situation
 - **A successor/result function (with actions, costs)**
 - Actions are choices an agent could take in each state
 - **A start state**
 - Where the game starts
 - **A goal test**
 - When the game ends


Real world task - Pac-man



Search problem formulation - Pac-man

- A search problem consists of:


- A state space



- A successor/result function (with actions, costs)


- Actions: Up, Down, Left, Right
- Cost: 1 for each step

- A start state



- A goal test

- Move to a specific position



Real world task : the 8-puzzle

You can move each tile to the neighbor position if it is blank

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

Search problem formulation : the 8-puzzle

7	2	4
5		6
8	3	1


Start State

1	2	3
4	5	6
7	8	

Goal State


- State space
 - Locations of 8 tiles. Each element contains a (x,y).
- Successor function:
 - Move blank (up, down, right, left)

Real Task: 8-Queens Puzzle



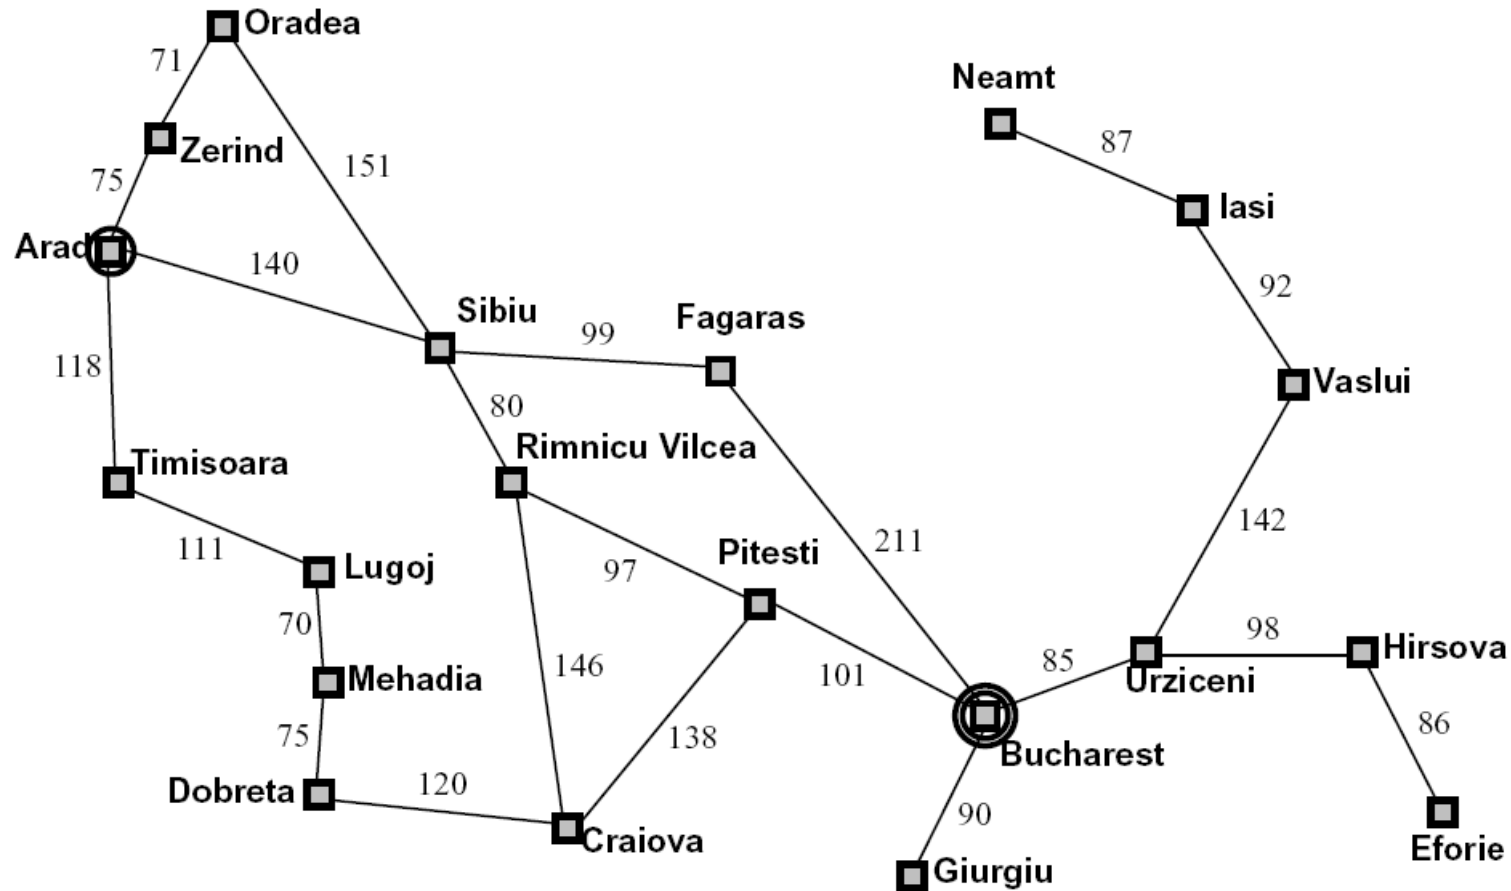
- Place 8 queens on a chessboard and ensure no two queens attack each other.
- A queen attacks any piece in the same row, column or diagonal.
- In this case, 3 more queens missing

Search problem formulation : 8-Queens Puzzle



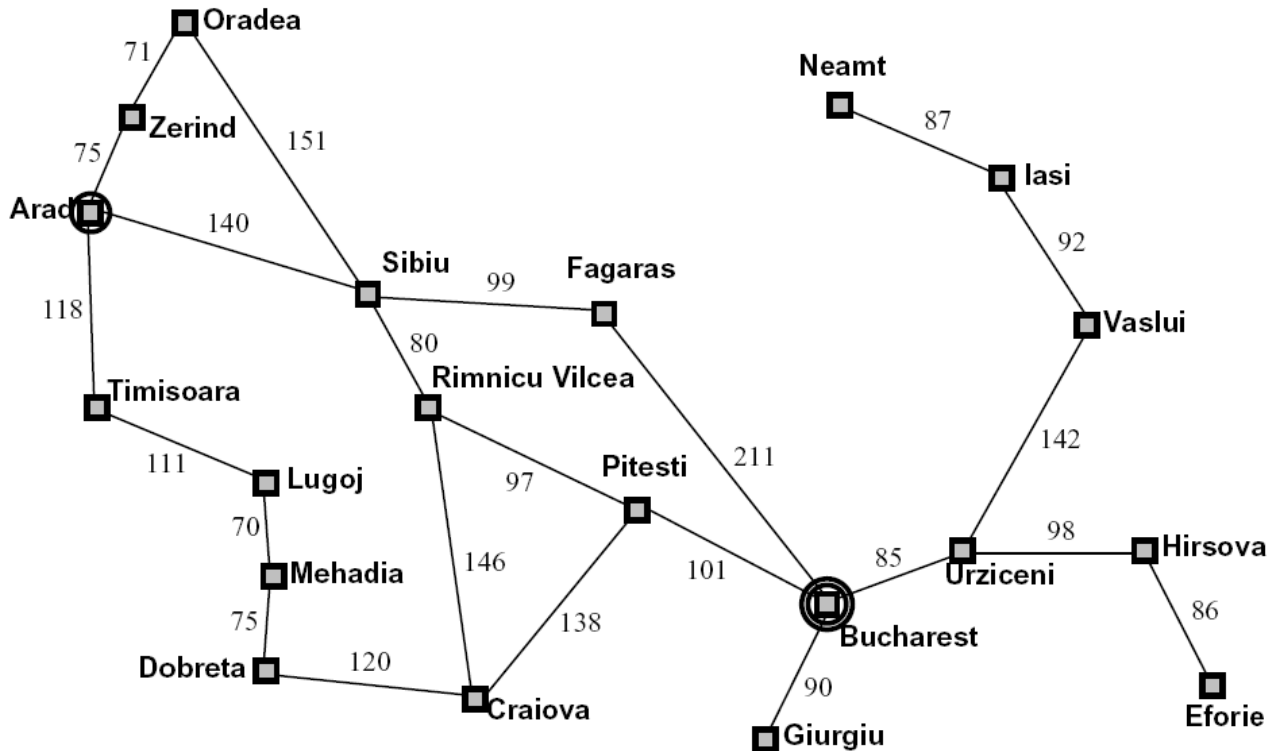
- **State space**
 - any arrangement of 0 to 8 queens on board. A vector a 8X8 elements, 0/1.
- **Successor function**
 - add a queen to any square
- **Start state**
 - blank board
- **Goal test**
 - 8 queens on board, non attacked

Real world task : Traveling in Romania



- Travel from one city to a target city.
- Each node stands for a city. Numbers on edges are distances.
- Identify a shortest path from origin city to target city.
- In this case from Arad to Bucharest.

Search problem formulation : Traveling in Romania




- State space
 - Cities
- Successor function:
 - Action: Go to adjacent city
 - Cost: distance
- Start state:
 - Arad
- Goal test:
 - Is state == Bucharest?

Outline

- **Search Problems**
- **State Spaces**

What is in a State Space?

The **world state** includes every last detail of the environment




A **search state** keeps only the details needed for solving a specific problem

- Problem: Path-Finding
 - States: (x,y) **location**
 - Actions: NSEW
 - Successor: update location following an action
 - Goal test: is $(x,y)=\text{END}$
- Problem: Eat-All-Dots
 - States: $\{ (x,y), \text{dot map} \}$
 - Actions: NSEW
 - Successor: update location and dot distribution
 - Goal test: dots all false


State Space Sizes?

- World state:
 - Board blanks: 64
 - Queen number: 8
- How many
 - World states?
 - 64^8



State Space Sizes?

- World state:
 - Agent positions: 120
 - Food count: 30
 - Ghost positions: 12
 - Agent facing: NSEW
- How many
 - World states?
 $120 \times (2^{30}) \times (12^2) \times 4$
 - States for path-finding?
120
 - States for eat-all-dots?
 $120 \times (2^{30})$



Outline


- **Search Problems**
- **State Spaces**
- **State Space Representation**

State Space Graph

- State space graph: A mathematical representation of a search problem


- **Nodes** are (abstracted) **states**
- **Arcs** represent **successors** (action results)
- The goal test is a **set of goal nodes** (maybe only one)

- In a state space graph, each state occurs only once.
- Nodes corresponds to states in the state space one to one.




Search Trees

- The start state is the root node
- Children correspond to successors
- Nodes show states, but correspond to **PLANS** that achieve those states



- **PLAN** means a series of actions.


Transforms this State Space Graph to a Search Tree



- *Each NODE in the search tree is an entire PATH in the state space graph.*
- *Each state can be corresponded to multiple nodes in the tree.*

State Space Graphs vs. Search Trees

Consider this 4-state graph:



How big is its search tree (from S)?

Important: Lots of repeated structure in the search tree!

For most problems, we can never actually build the whole tree.

Outline

- **Search Problems**
- **State Spaces**
- **State Space Representation**
- **Solve a Search Problem**

Solving a Search Problem

- A search problem consists of:
 - A state space
 - A successor/result function (with actions, costs)
 - A start state
 - A goal test
- Problem Representation
 - State Space Graph
 - State Search Tree
- Solving a search problem is to identify a sequence of actions from start state to goal state.
- Also known as planning

Standard Search Setup

- Process
 - **Start** at the initial state from the search tree
 - **Expand** an un-visited but observed node
 - **Generate** neighbors of current state (observe new nodes)
 - **Test** whether the state matches the goal configuration
- Implementation Components
 - **Fringe**: store those nodes expanded but not visited
 - Expand: numerate neighbors of a node
 - Generate: add a node into fringe
 - **Search strategy**: choose a node from fringe to visit

Pseudo-Code for Tree Search Paradigm


```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe ← INSERT(child-node, fringe)
    end
  end
end
```

- **Fringe**: store those nodes expanded but not visited
 - Expand: EXPAND
 - Generate: INSERT
- **Search strategy**: REMOVE-FRONT

Goal test is performed when the node is removed from the fringe.

Evaluation of Search Algorithm

- Some symbolic of search tree:
 - b is the branching factor
 - m is the maximum depth
 - s are the solutions at various depths



- Number of nodes in entire tree?

$$1 + b + b^2 + \dots + b^m = \frac{1-b^{m+1}}{1-b} = O(b^m)$$

- **Complete:** guaranteed to find a solution if one exists?
- **Optimal:** guaranteed to find the least cost path?
- **Time complexity:** # of nodes generated
- **Space complexity:** # of nodes stored (size of fringe)

Outline

- **Search Problems**
- **State Spaces**
- **State Space Representation**
- **Solve a Search Problem**
- **Uninformed Search Algorithms**
 - **Depth-First Search**
 - **Breadth-First Search**
 - **Iterative Deepening**
 - **Uniform-Cost Search**

Depth-First Search



Fringe: A data structure to store nodes observed but not visited

Expand: pop a node from the fringe

Generate: get neighbors of the expanded node and insert them into the fringe


Strategy: expand a deepest node first (number of steps to the root)

Implementation: Fringe is a LIFO stack



Depth-First Search (DFS) Properties

- How many nodes does DFS generate (Time Complexity)?
 - If m is finite, takes time $O(b^m)$
- How much space does the fringe take (Space Complexity)?
 - b nodes in each layer, $O(bm)$
- Is it complete?
 - No, m can be infinite
- Is it optimal?
 - No, it finds the “leftmost” solution, regardless of depth or cost



Breadth-First Search (BFS)


Fringe: A data structure to store nodes observed but not visited

Expand: pop a node from the fringe


Generate: get neighbors of the expanded node and insert them into the fringe

Strategy: expand a shallowest node first (number of steps to the root)

Implementation: Fringe is a FIFO queue




Search
Tiers



Breadth-First Search (BFS) Properties

- How many nodes does BFS expand (Time Complexity)?
 - Processes all nodes above shallowest solution
 - Let depth of shallowest solution be s
 - Search takes time $O(b^{s+1})$. Why $s+1$?
- How much space does the fringe take?
 - Has roughly the last tier, so $O(b^{s+1})$
- Is it complete?
 - yes
- Is it optimal?
 - Yes (if the cost is equal per step)



DFS vs BFS

- DFS requires smaller storage
- BFS is robust (complete and optimal) but needs more space for storage


How bad is BFS?

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes


Figure 3.13 Time and memory requirements for breadth-first search. The numbers shown assume branching factor $b = 10$; 1 million nodes/second; 1000 bytes/node.

Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
 - Run a DFS with depth limit 1. If no solution...
 - Run a DFS with depth limit 2. If no solution...
 - Run a DFS with depth limit 3.
- How many nodes does BFS generate?
 - $O(b^{s+1})$
- How much space does the fringe take?
 - $O(bs)$
- Although the time complexity of IDS is in the same order as BFS, it generates many more nodes.
 - $b + (b + b^2) + (b + b^2 + b^3) \dots$
 - $= sb^1 + (s-1)b^2 + (s-2)b^3 + (s-3)b^4 \dots + b^{s+1}$
 - $= O(sb^1 + (s-1)b^2 + (s-2)b^3 + (s-3)b^4 \dots + b^{s+1}) = O(b^{s+1})$



Cost-Sensitive Search



- BFS finds the shortest path in terms of number of actions. It does not find the least-cost path.
- We will now cover a similar algorithm which does find the least-cost path.

Uniform Cost Search



Fringe: A data structure to store nodes observed but not visited

Expand: pop a node from the fringe

Generate: get neighbors of the expanded node and insert them into the fringe


Strategy: expand a node with least cost (accumulative cost from root to this node)

Implementation: Fringe is a Priority queue



Uniform Cost Search (UCS) Properties

- Time Complexity
 - Takes time $O(b^{C^*/\epsilon+1})$ (exponential in effective depth)
- Space Complexity
 - Has roughly the last tier, so $O(b^{C^*/\epsilon+1})$
- Is it complete?
 - yes!
 - best solution has a finite cost
 - minimum arc cost is positive.
- Is it optimal?
 - Yes! minimum arc cost is positive.



C^* : cost for the lowest solution

ϵ : the cost of least arc in the search tree

C^*/ϵ : effective depth, most steps from root to the solution

Comparison

Algorithm	Complete?	Optimal?	Time?	Space?
DFS	N	N	$O(b^m)$	$O(bm)$
BFS	Y	Y	$O(b^{s+1})$	$O(b^{s+1})$
IDS	Y	Y	$O(b^{s+1})$	$O(bs)$
UCS	Y	Y	$O(b^{C^*/\epsilon+1})$	$O(b^{C^*/\epsilon+1})$

b is branching factor

s is the depth of the optimal solution

m is the maximum depth of the tree

C^* is cost for the lowest solution

ϵ is the cost of least arc in the search tree



For BFS, Suppose the branching factor b is finite and step costs are identical

Outline

- **Search Problems**
- **State Spaces**
- **State Space Representation**
- **Solve a Search Problem**
- **Uninformed Search Algorithms**
 - **Depth-First Search**
 - **Breadth-First Search**
 - **Iterative Deepening**
 - **Uniform-Cost Search**
- **Tree Search VS Graph Search**


Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work.



Motivation of Graph Search

- We shouldn't bother expanding the circled nodes (why?)
- We have visited their related states earlier with less cost



Graph Search

- Idea: never **expand** a state twice
- Tree search + set of expanded states (“closed set”)
- Update the expanded states set after expanding a state
- Check the expanded states set before expanding a state, if it has been added into the set, skip it.

Pseudo-Code for Search Paradigm

```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe ← INSERT(child-node, fringe)
    end
  end
```

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
    end
  end
```